



Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:04/09/23
Date of Submission: 27/09/23



Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum_j w_j * \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for



To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0
2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.



capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

1. The Gradient Boosting Classifier demonstrates the best performance on this dataset, boasting the highest accuracy (0.8732) and F1 score (0.71).
2. They exhibit commendable precision levels, implying their ability to accurately predict the positive class (1).
3. Among the classifiers, the Gradient Boosting Classifier stands out with the highest recall for the positive class (1), signifying its proficiency in correctly identifying more positive cases.
4. The F1 score, which strikes a balance between precision and recall, effectively gauges how well these classifiers perform in classifying both classes.

Comparison between Boosting Algorithm and Random Forest Classifier applied on the Adult Income Dataset :

1. Boosting algorithms, including Gradient Boosting, and XGBoost, consistently exhibit superior performance compared to the Random Forest Classifier, boasting higher accuracy, precision, and F1 scores.
2. The Random Forest Classifier maintains a respectable level of performance, achieving an accuracy rate of approximately 85% and a well-balanced F1 score. Nevertheless, it lags behind the boosting algorithms in terms of overall performance.
3. Boosting algorithms consistently demonstrate greater precision and recall for the positive class (income > 50K), signifying their enhanced capability to accurately classify individuals with high incomes.



4. Across the board, all these models deliver commendable results when applied to the Adult Census Income Dataset, with the boosting algorithms, particularly Gradient Boosting, showcasing slightly superior accuracy and F1 score compared to the Random Forest Classifier.

Random Forest Accuracy: 85.40

Boosting algorithm Accuracy: 87.05

```
# importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

df = pd.read_csv('/content/adult.csv')
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	ca
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female		0
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female		0
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female		0
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female		0
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female		0

```
df.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

▾ Data Treatment

```
df.isna().sum()

age          0
workclass    0
fnlwgt       0
education    0
education.num 0
marital.status 0
occupation   0
relationship 0
race         0
sex          0
capital.gain  0
capital.loss  0
hours.per.week 0
native.country 0
income       0
dtype: int64

df.loc[df.duplicated() == True]
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain
8453	25	Private	308144	Bachelors	13	Never-married	Craft-repair	Not-in-family	White	Male	0
8645	90	Private	52386	Some-college	10	Never-married	Other-service	Not-in-family	Asian-Pac-Islander	Male	0
12202	21	Private	250051	Some-college	10	Never-married	Prof-specialty	Own-child	White	Female	0
14346	20	Private	107658	Some-college	10	Never-married	Tech-support	Not-in-family	White	Female	0
15603	25	Private	195994	1st-4th	2	Never-married	Priv-house-serv	Not-in-family	White	Female	0
17344	21	Private	243368	Preschool	1	Never-married	Farming-fishing	Not-in-family	White	Male	0
19067	46	Private	173243	HS-grad	9	Married-civ-spouse	Craft-repair	Husband	White	Male	0
20388	30	Private	144593	HS-grad	9	Never-married	Other-service	Not-in-family	Black	Male	0
20507	19	Private	97261	HS-grad	9	Never-married	Farming-fishing	Not-in-family	White	Male	0
22783	19	Private	138153	Some-college	10	Never-married	Adm-clerical	Own-child	White	Female	0
22934	19	Private	146679	Some-college	10	Never-married	Exec-managerial	Own-child	Black	Male	0
23276	49	Private	31267	7th-8th	4	Married-civ-spouse	Craft-repair	Husband	White	Male	0
23660	25	Private	195994	1st-4th	2	Never-married	Priv-house-serv	Not-in-family	White	Female	0
23720	44	Private	367749	Bachelors	13	Never-married	Prof-specialty	Not-in-family	White	Female	0
23827	49	Self-emp-not-inc	43479	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	Male	0
26738	23	Private	240137	5th-6th	3	Never-married	Handlers-cleaners	Not-in-family	White	Male	0
27133	28	Private	274679	Masters	14	Never-married	Prof-specialty	Not-in-family	White	Male	0
28796	27	Private	255582	HS-grad	9	Never-married	Machine-op-inspct	Not-in-family	White	Female	0
29054	42	Private	204235	Some-college	10	Married-civ-spouse	Prof-specialty	Husband	White	Male	0
29334	39	Private	30916	HS-grad	9	Married-civ-spouse	Craft-repair	Husband	White	Male	0

df = df.drop_duplicates()

df.loc[df.duplicated() == True]

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss
31060	46	Private	133616	HS-grad	10	Divorced	Unemployed	Unmarried	White	Female	0	0

df['age'].describe()

count

32537.000000

mean

38.585549

std

13.637984

min

17.000000

25%

28.000000

50%

37.000000

75%

48.000000

max

90.000000


Name: age, dtype: float64

➤ Exploratory Data Analysis

Analysis per variable

```
# Age
plt.figure(figsize=(16, 8))
sns.set_theme(style="darkgrid")
sns.set_palette("magma")
sns.histplot(data=df, x='age', hue='income', bins=20, multiple='stack')
```

```
plt.xlabel('Age')
plt.ylabel('Count')
plt.title('Age Distribution')
plt.show()
```



```
# workclass
df.groupby('workclass').size()
```

```
workclass
?          1836
Federal-gov  960
Local-gov   2093
Never-worked    7
Private     22673
Self-emp-inc  1116
Self-emp-not-inc 2540
State-gov    1298
Without-pay    14
dtype: int64
```

```
workclass_unknown = df.loc[df['workclass'] == '?']
print('**age distribution for workclass "?"** \n', workclass_unknown['age'].describe())
plt.figure(figsize=(16, 8))
plt.title('age distribution for workclass "?"')
plt.hist(workclass_unknown['age'], bins=20)
sns.histplot(data=df.loc[df['workclass'] == '?'], x='age', hue='income', bins=20, multiple='stack')
```



```
**age distribution for workclass "?"**
```

```
count    1836.000000
mean      40.960240
std       20.334587
min       17.000000
25%       21.000000
50%       35.000000
75%       61.000000
max       90.000000
```

```
Name: age, dtype: float64
```

```
<Axes: title={'center': 'age distribution for workclass "?"'}, xlabel='age', ylabel='Count'>
```



```
print(df.query('age < 20').groupby('workclass').size())
print(df.query('age > 20 and age < 60').groupby('workclass').size())
print(df.query('age > 60').groupby('workclass').size())
```

```
workclass
?          269
Federal-gov    9
Local-gov     35
Never-worked   4
Private      1249
Self-emp-inc   16
Self-emp-not-inc 37
State-gov     32
Without-pay    2
dtype: int64

workclass
?          928
Federal-gov   874
Local-gov    1875
Never-worked   2
Private     19599
Self-emp-inc   940
Self-emp-not-inc 2108
State-gov    1158
Without-pay    5
dtype: int64

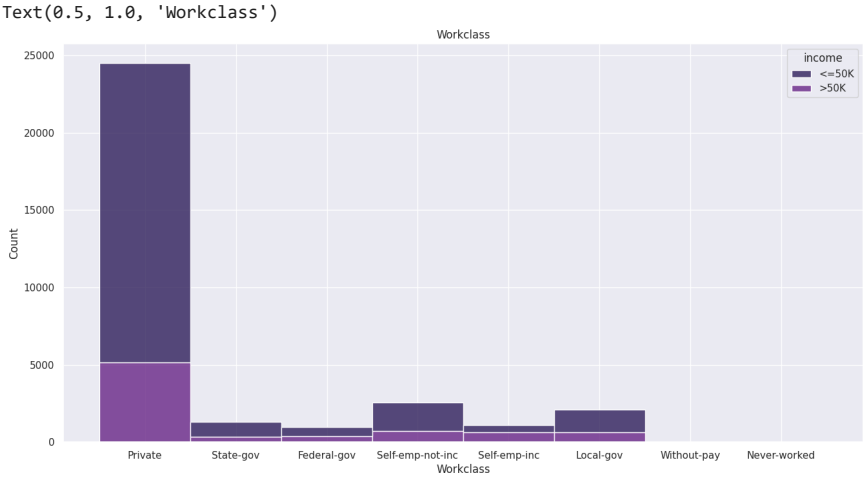
workclass
?          493
Federal-gov    58
Local-gov     151
Private     1070
Self-emp-inc   143
Self-emp-not-inc 338
State-gov     71
Without-pay    7
dtype: int64
```

```
df.loc[df['workclass'] == '?', 'workclass'] = 'Private'
df.loc[df['workclass'] == '?']
```

```
age workclass fnlwgt education education.num marital.status occupation relationship race sex capital.gain capital.loss
```

```
plt.figure(figsize=(16, 8))
sns.histplot(data=df, x='workclass', hue='income', multiple='stack')
plt.xlabel('Workclass')
plt.title('Workclass')
```





```
df.groupby(df['workclass']).size()
```


```
workclass
Federal-gov      960
Local-gov       2093
Never-worked        7
Private       24509
Self-emp-inc    1116
Self-emp-not-inc 2540
State-gov      1298
Without-pay      14
dtype: int64
```

```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.l
0	90	Private	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	4
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	4
2	66	Private	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	4

Machine-

```
# Final Weight
plt.figure(figsize=(16, 8))
sns.histplot(x='fnlwgt', data=df, hue='income', multiple='stack')
plt.title('Final Weight distribution')
plt.xlabel('Final Weight')
```



```
# Education
df.groupby('education').size()
```

education	
10th	933
11th	1175
12th	433
1st-4th	166
5th-6th	332
7th-8th	645
9th	514
Assoc-acdm	1067
Assoc-voc	1382
Bachelors	5353
Doctorate	413
HS-grad	10494
Masters	1722
Preschool	50
Prof-school	576
Some-college	7282
dtype: int64	

```
plt.figure(figsize=(16, 8))
plt.pie(df.groupby('education').size(), labels=df.groupby('education').size().index, autopct='%1.1f%%')
```

```

Text(-0.7785868926457222, 0.7770472640710339, 'Doctorate'),
Text(-1.0605252834929717, -0.2920378795159155, 'HS-grad'),
Text(-0.13449613832544505, -1.091746668772359, 'Masters'),
Text(0.05334950660702595, -1.0987055247630217, 'Preschool'),
Text(0.11962081577950918, -1.0934765020027841, 'Prof-school'),
Text(0.8391187061846062, -0.7112522737616188, 'Some-college']],
[Text(0.5975670366161832, 0.05397811361795755, '2.9%'),
Text(0.5743211004298473, 0.17365273853599109, '3.6%'),
Text(0.540559772491962, 0.2603749841352626, '1.3%'),
Text(0.5246052478260459, 0.29118608131806883, '0.5%'),
Text(0.510002867060924, 0.3160649863392614, '1.0%'),
Text(0.47796394868079695, 0.3626988609872665, '2.0%'),
Text(0.4344705359375361, 0.4138059368860603, '1.6%'),
Text(0.36649465623567234, 0.4750596456769363, '3.3%'),
Text(0.2450065934522299, 0.547696785790216, '4.2%'),
Text(-0.13658484476458072, 0.5842470198303414, '16.5%'),
Text(-0.42468375962493937, 0.42384396222056386, '1.3%'),
Text(-0.5784683364507117, -0.15929338882686297, '32.3%'),
Text(-0.0733615299956973, -0.5954981829667412, '5.3%'),
Text(0.029099730876559603, -0.5992939225980118, '0.2%'),
Text(0.0652477176979141, -0.596441728365155, '1.8%'),
Text(0.4577011124643306, -0.38795578568815564, '22.4%')]]


```

```

plt.figure(figsize=(16, 8))
sns.countplot(x='education', data=df, hue='income')

```

<Axes: xlabel='education', ylabel='count'>



```
# Education Number
```

```
df.groupby('education.num').size()
```

```
education.num
```

```
1      50
```

```
2     166
```

```
3     332
```

```
4     645
```

```
5     514
```

```
6     933
```

```
7    1175
```

```
8     433
```

```
9   10494
```

```
10   7282
```

```
11   1382
```

```
12   1067
```

```
13   5353
```

```
14   1722
```

```
15     576
```

```
16     413
```

```
dtype: int64
```

```
df['education.num'].describe()
```

```
count    32537.000000
```

```
mean       10.081815
```

```
std         2.571633
```

```
min         1.000000
```

```
25%         9.000000
```

```
50%        10.000000
```

```
75%        12.000000
```


```
max        16.000000
```

```
Name: education.num, dtype: float64
```

```
plt.figure(figsize=(16, 8))
```

```
sns.countplot(x='education.num', data=df, hue='income')
```

<Axes: xlabel='education.num', ylabel='count'>




```
# marital status
```

```
print(df.groupby('marital.status').size())
```


```
plt.figure(figsize=(16, 8))
sns.countplot(data=df, x='marital.status', hue='income')
```

```
marital.status
Divorced          4441
Married-AF-spouse    23
Married-civ-spouse 14970
Married-spouse-absent  418
Never-married     10667
Separated         1025
Widowed           993
dtype: int64
<Axes: xlabel='marital.status', ylabel='count'>
```

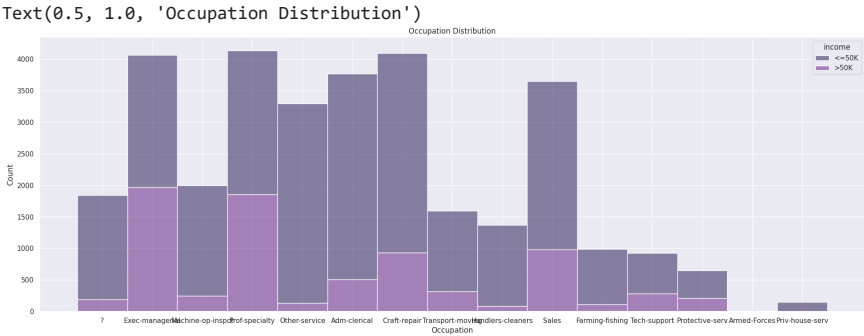


```
# Occupation
print(df.groupby('occupation').size())
plt.figure(figsize=(16, 12))
plt.pie(df.groupby('occupation').size(), labels=df.groupby('occupation').size().index, autopct='%1.1f%%')
plt.title('Occupation Distribution')
```

```
occupation
?          1843
Adm-clerical 3768
Armed-Forces 9
Craft-repair 4094
Exec-managerial 4065
Farming-fishing 992
Handlers-cleaners 1369
Machine-op-inspct 2000
Other-service 3291
Priv-house-serv 147
Prof-specialty 4136
Protective-serv 649
Sales 3650
Tech-support 927
Transport-moving 1597
dtype: int64
Text(0.5, 1.0, 'Occupation Distribution')
```



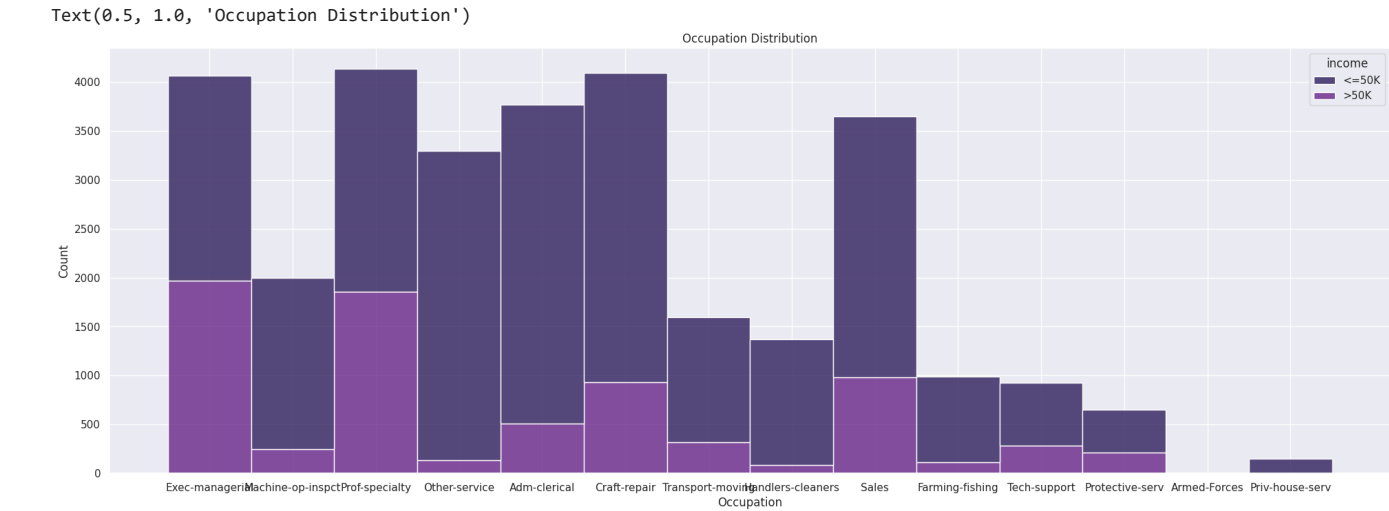
```
plt.figure(figsize=(24, 8))
sns.histplot(data=df, x='occupation', hue='income', multiple='stack', alpha=0.5)
plt.xlabel('Occupation')
plt.title('Occupation Distribution')
```




```
df.drop(df.loc[df['occupation'] == '?'].index, inplace=True)

plt.figure(figsize=(24, 8))
sns.histplot(data=df, x='occupation', hue='income', multiple='stack')
```

```
plt.xlabel('Occupation')
plt.title('Occupation Distribution')
```



```
# Relationship
print(df.groupby('relationship').size())
plt.figure(figsize=(16, 8))
sns.histplot(data=df, x='relationship', hue='income', multiple='stack')
```

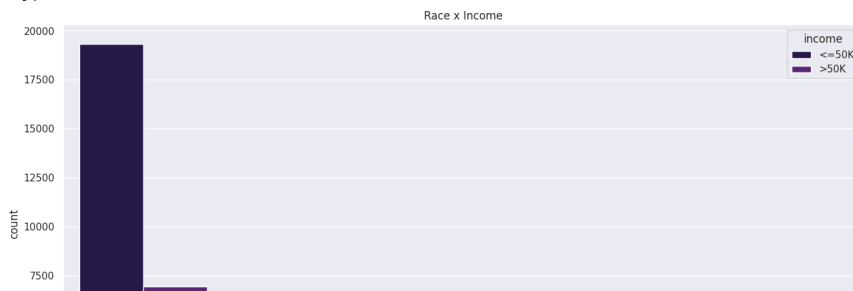



```
# Race and Sex
plt.figure(figsize=(16, 8))
print(df.groupby(df.race).size())
sns.countplot(data=df, x='race', hue='income')
plt.title('Race x Income')
plt.show()
plt.figure(figsize=(16, 8))
print(df.groupby(df.sex).size())
sns.countplot(data=df, x='sex', hue='income')
plt.title('Sex x Income')
```

```

race
Amer-Indian-Eskimo    286
Asian-Pac-Islander   973
Black                 2907
Other                 248
White                26280
dtype: int64

```



```

# Capital Gain & Capital Loss
print('*** capital gain *** \n ', df.groupby('capital.gain').size(), '\n')
print('*** capital loss *** \n ', df.groupby('capital.loss').size(), '\n')

```

```

*** capital gain ***
capital.gain
0      28105
114      6
401      1
594     29
914      8
...
25236    11
27828    33
34095     3
41310     2
99999    155
Length: 118, dtype: int64

```

```

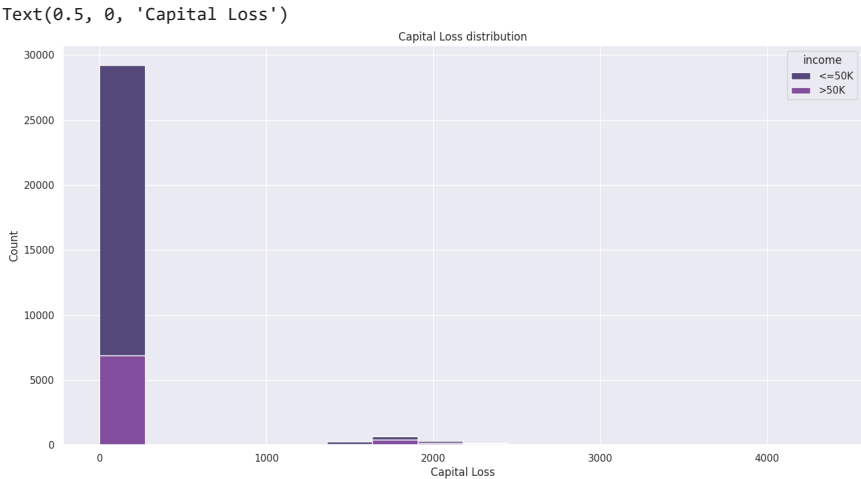
*** capital loss ***
capital.loss
0      29233
155      1
213      4
323      3
419      1
...
3004     2
3683     2
3770     2
3900     2
4356     1
Length: 90, dtype: int64

```

```

plt.figure(figsize=(16, 8))
sns.histplot(x='capital.gain', data=df, hue='income', multiple='stack')
plt.title('Capital Gain distribution')
plt.xlabel('Capital Gain')
plt.show()
plt.figure(figsize=(16, 8))
sns.histplot(x='capital.loss', data=df, hue='income', multiple='stack')
plt.title('Capital Loss distribution')
plt.xlabel('Capital Loss')

```



Text(0.5, 0, 'Capital Loss')

```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	rela
1	82	Private	132870	HS-grad	9	Widowed	Exec-manual	No
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	l
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	

```
# Hours per week
df.groupby('hours.per.week').size()
```

hours.per.week	
1	8
2	15
3	24
4	28
5	39
..	
95	2
96	5

```

97      2
98     11
99     80
Length: 94, dtype: int64

```

```
df = df.drop(columns=['hours.per.week'])
```

```

# Native Country
df.groupby('native.country').size()

```

```

native.country
?                    555
Cambodia             18
Canada              107
China                68
Columbia             56
Cuba                 92
Dominican-Republic  67
Ecuador             27
El-Salvador         100
England              86
France              27
Germany             128
Greece              29
Guatemala           61
Haiti                42
Holand-Netherlands   1
Honduras            12
Hong                19
Hungary             13
India               100
Iran                42
Ireland             24
Italy               68
Jamaica             80
Japan               59
Laos                17
Mexico              606
Nicaragua           33
Outlying-US(Guam-USVI-etc) 14
Peru                30
Philippines         188
Poland              56
Portugal            34
Puerto-Rico        109
Scotland            11
South               71
Taiwan              42
Thailand            17
Trinidad&Tobago     18
United-States       27487
Vietnam             64
Yugoslavia          16
dtype: int64

```

```

plt.figure(figsize=(16, 8))
plt.pie(df.groupby('native.country').size(), labels=df.groupby('native.country').size().index, autopct='%1.1f%%')

```

▼ Model preparation & building

```

label_encoder = LabelEncoder()
categorical_columns = ['income', 'workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
df[categorical_columns] = df[categorical_columns].apply(label_encoder.fit_transform)

```

```
df
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
1	82	2	132870	11	9	6	3
3	54	2	140359	5	4	0	6
4	41	2	264663	15	10	5	9
5	34	2	216864	11	9	0	7

```
x_train, x_test, y_train, y_test = train_test_split(df.drop(columns=['income']), df['income'], test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

Random Forest

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

rfc = RandomForestClassifier()
rfc.fit(x_train, y_train)
y_pred = rfc.predict(x_test)
print('**** ACCURACY_SCORE **** \n\n', accuracy_score(y_test, y_pred), '\n')
print('**** CLASSIFICATION_REPORT **** \n\n', classification_report(y_test, y_pred), '\n')
print('**** CONFUSION MATRIX ****')
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
```

**** ACCURACY_SCORE ****


0.8540478905359179

**** CLASSIFICATION_REPORT ****

	precision	recall	f1-score	support
0	0.89	0.93	0.91	4699
1	0.72	0.62	0.66	1440
accuracy			0.85	6139
macro avg	0.80	0.77	0.79	6139
weighted avg	0.85	0.85	0.85	6139

**** CONFUSION MATRIX ****

<Axes: >



XGBoost Classifier

```
from xgboost import XGBClassifier
```

```
xgb = XGBClassifier()
xgb.fit(x_train, y_train)
y_pred = xgb.predict(x_test)
```

```
print('**** ACCURACY_SCORE **** \n\n', accuracy_score(y_test, y_pred), '\n')
print('**** CLASSIFICATION_REPORT **** \n\n', classification_report(y_test, y_pred), '\n')
print('**** CONFUSION MATRIX ****')
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
```

**** ACCURACY_SCORE ****


0.8705000814464896

**** CLASSIFICATION_REPORT ****

	precision	recall	f1-score	support
0	0.90	0.94	0.92	4699
1	0.76	0.65	0.70	1440
accuracy			0.87	6139
macro avg	0.83	0.79	0.81	6139
weighted avg	0.87	0.87	0.87	6139

**** CONFUSION MATRIX ****

<Axes: >



```
from sklearn.ensemble import GradientBoostingClassifier
```

#Training the model with gradient boosting

```
gbc = GradientBoostingClassifier(
    learning_rate=0.1,
    n_estimators = 500,
    max_depth=5,
    subsample=0.9,
    min_samples_split = 100,
    max_features='sqrt',
    random_state=10)
gbc.fit(x_train,y_train)

# Predictions
y_pred_gbc =gbc.predict (x_test)
print("Accuracy: ",accuracy_score (y_test, y_pred_gbc))
print('**** CLASSIFICATION_REPORT **** \n\n', classification_report(y_test, y_pred_gbc), '\n')
print('**** CONFUSION MATRIX ****')
sns.heatmap(confusion_matrix(y_test, y_pred_gbc), annot=True, fmt='d')
```